University of Toronto, Department of Computer Science
CSC 485H/2501H: Computational linguistics, Fall 2022

程序代写代做CS编程辅导

Assignment 2

Due date: 23:59 on F[illegible] 2022.
Late assignments wil[illegible]out a valid medical certificate or other documentation of an emergency.
For CSC485 students[illegible]orth 33% of your final grade.
For CSC2501 students, this assignment is worth 25% of your final grade.

WeChat: cstutorcs

- **Read the whole assignment carefully.**
- Type the written parts of your submission in no less than 12pt font.
- What you turn in must be your own work. You may not work with anyone else on any of the problems in this assignment. If you need assistance, contact the instructor or TA for the assignment.

Assignment Project Exam Help

- Any clarifications to the problems will be posted on the Piazza forum for the class. You will be responsible for taking into account in your solutions any information that is posted there, or discussed in class, so you should check the page regularly between now and the due date.

Email: tutorcs@163.com

- The starter code directory for this assignment is accessible on Teaching Labs machines at the path `/u/csc485h/fall/pub/a2/`. In this handout, code files we refer to are located in that directory.

QQ: 749389476

- When implementing code, make sure to **read the docstrings** as some of them provide important instructions, implementation details, or hints.

https://tutorcs.com

- Fill in your name, student number, and UTORid on the relevant lines at the top of each file that you submit. (Do not add new lines; just replace the `NAME`, `NUMBER`, and `UTORid` placeholders.)

程序代写代做 CS编程辅导

0. Warming up with WordNet and NLTK (4 marks)

WordNet is a lex dictionary, WordNet provides definitions and example usages for different s s. But WordNet does the job of a thesaurus as well: it provides synonyms fo g synonymous senses together into a set called a *synset*.

But wait; there's provides information about semantic relationships beyond synonymy, such onymy/hyponymy, and meronymy/holonymy. Throughout this assignment, y e of WordNet via the NLTK package, so the first step is to get acquainted with doing so. Consult sections 4.1 and 5 of chapter 2 as well as section 3.1 of chapter 3 of the NLTK book for an introduction along with examples that you will likely find useful for this assignment. You may also find section 3.6 is also useful for its discussion of lemmatization, although you will not be doing any lemmatization for this assignment.

Make certain that you use Python 3 (`python3`). That is where NLTK lives on our machines. You will need to `import nltk`. You may also need to `nltk.download('omw-1.4')`.

WeChat: cstutorcs
Assignment Project Exam Help
Email: tutorcs@163.com
QQ: 749389476
https://tutorcs.com

(a) (1 mark) A root hyperonym is a synset with no hyperonyms. A synset s is said to have depth d if there are d hyperonym links between s and a root hyperonym. Keep in mind that, because synsets can have multiple hyperonyms, they can have multiple paths to root hyperonyms.

Implement the `deepest` function in `q0.py` that finds the synset in WordNet with the largest maximum depth and report both the synset and its depth on each of its paths to a root hyperonym.[1]

(b) (2 marks) Implement the `superdefn` function in `q0.py` that takes a synset `s` and returns a list consisting of all of the tokens in the definitions of `s`, its hyperonyms, and its hyponyms. Use `word_tokenize` as shown in chapter 3 of the NLTK book.

(c) (1 mark) WordNet's `word_tokenize` only tokenizes text; it doesn't filter out any of the tokens. You will be calculating overlaps between sets of strings, so it will be important to remove stop words and any tokens that consist entirely of punctuation symbols.

Implement the `stop_tokenize` function in `q0.py` that takes a string, tokenizes it using `word_tokenize`, removes any tokens that occur in NLTK's list of English stop words (which has already been imported for you), and also removes any tokens that consist entirely of punctuation characters. For a list of punctuation symbols, use Python's `punctuation` characters from the `string` module (this has also already been imported for you). Keep in mind that NLTK's list contains only lower-case tokens, but the input string to `stop_tokenize` may contain upper-case symbols. Maintain the original case in what you return.

[1] Hint: you may find the `wn.all_synsets` and `synset.max_depth` methods helpful.

程序代写代做 CS编程辅导

1. The Lesk algorithm & word2vec (28 marks)

Recall the proble ambiguation (WSD): given a semantically ambiguous word in context, det ense. A simple but surprisingly hard-to-beat baseline method for WSD is e (MFS): just select the most frequent sense for each ambiguous word, wh are provided by some corpus.

(a) (1 mark) Imple n that returns the most frequent sense for a given word in a sentence. Note that `wordnet.synsets()` orders its synsets by decreasing frequency.

As discussed in class, the Lesk algorithm is a venerable method for WSD. The Lesk algorithm variant that we will be using for this assignment selects the sense with the largest largest number of words in common with the ambiguous word's sentence. This version is called the simplified Lesk algorithm.

WeChat: cstutorcs
Assignment Project Exam Help

Algorithm 1: The simplified Lesk algorithm.

input : a word to disambiguate and the sentence in which it appears

```
best_sense ⟵ most_frequent_sense(word)
best_score ⟵ 0
context ⟵ the bag of words in sentence
for each sense of word do
    signature ⟵ the bag of words in the definition and examples of sense
    score ⟵ Overlap(signature, context)
    if score > best_score then
        best_sense ⟵ sense
        best_score ⟵ score
    end
end
return best_sense
```

Email: tutorcs@163.com
QQ: 749389476
https://tutorcs.com

Our version represents the signature and context each as *bags* (also known as *multisets*) of words. Bags are like sets but allow for repeated elements by assigning each element with a non-negative integer that indicates the number of instances of that element in the bag; this integer is called *multiplicity*. Because of multiplicity, the bags $A = \{a,a,b\}$, $B = \{a,b,b\}$, $C = \{a,b\}$, and $D = \{a,a,b,b\}$ are all different. (This is not the case for sets, which do not allow multiplicity.) a has multiplicity 2 in A and D but multiplicity 1 in B and C, while b has multiplicity 2 in B and D but multiplicity 1 in A and C.

As with sets, each bag has associated with it a *cardinality* which is the sum of the multiplicities of its elements; so A and B have cardinality 3 while C has cardinality 2 and D has cardinality 4. The intersection of two bags Y and Z is the bag where, the multiplicity of any element x is defined as the minimum of the multiplicities of x in Y and in Z. For the bags defined above, C is the intersection

of A and B. The union is analogously defined using maximum instead of minimum; the bag D is the union of A and B.

程序代写代做 CS编程辅导

(b) (6 marks) In the `lesk` function, implement the simplified Lesk algorithm as specified in Algorithm 1, in `verlap(signature, context)` returns the cardinality of the intersect `ature` and `context`, i.e., the number of words tokens that the signatu `n` common.

Use your `stop_` `o` tokenize the examples and definitions.

Next, we're going `ied` Lesk algorithm so that the sense signatures are more informative.

(c) (3 marks) In the `lesk_ext` function, implement a version of Algorithm 1 where, in addition to including the words in `sense`'s definition and examples, `signature` also includes the words in the definition and examples of `sense`'s hyponyms, holonyms, and meronyms. Beware that NLTK has separate methods to access member, part, and substance holonyms/meronyms; use all of them.

Use `stop_tokenize` as you did for `lesk`.

WeChat: cstutorcs
Assignment Project Exam Help

(d) (2 mark) This extension should yield improvement in the algorithm's accuracy. Why is this extension helpful? Justify your answer.[2]

Email: tutorcs@163.com

Beyond `Overlap`,there are other scores we could use. Recall *cosine similarity* from the lectures: for vectors $\vec{v}$ and $\vec{w}$ with angle θ between them, the cosine similarity `CosSim` is defined as:

QQ: 749389476

$$\texttt{CosSim}(\vec{v},\vec{w}) = \cos\theta = \frac{\vec{v}\cdot\vec{w}}{|\vec{v}||\vec{w}|}$$

Cosine similarity can be applied to any two vectors in the same space. In the Lesk algorithm, we compare contexts with sense signatures, both of which are bags of words. If, instead of bags, we produced vectors from the relevant sources (i.e., the words in the sentence for the contexts and the words in the relevant definitions and examples for the sense signatures), we could then use cosine similarity to score the two.

https://tutorcs.com

Perhaps the simplest technique for constructing vectors from bags of words is to assign one vector dimension for every word, setting the value for each dimension to the number of occurrences of the associated word in the bag. So $\{a,a,b\}$ might be represented with the vector $\begin{bmatrix}2 & 1\end{bmatrix}$ and $\{a,b\}$ with $\begin{bmatrix}1 & 1\end{bmatrix}$. If we were comparing $\{\texttt{new},\texttt{buffalo},\texttt{york}\}$ and $\{\texttt{buffalo},\texttt{buffalo},\texttt{like}\}$, we might use $\begin{bmatrix}1 & 0 & 1 & 1\end{bmatrix}$ and $\begin{bmatrix}2 & 1 & 0 & 0\end{bmatrix}$, respectively.

(e) (4 marks) In the `lesk_cos` function, implement a variant of your `lesk_ext` function that uses `CosSim` instead of `Overlap`. You will have to modify `signature` and `context` so that they are vector-valued; construct the vectors from the relevant tokens for each in the manner described above.

(Again, use `stop_tokenize` to get the tokens for the signature.)

[2] Hint: consider the likely sizes of the overlaps.

(f) (2 marks) In the `lesk_cos_oneside` function, implement a variant of your `lesk_cos` function that, when constructing the vectors for the signature and context, does not include words that occur only in the signature. For example, if the signature has words {`new`, `buffalo`, `york`} while the context has {`buffalo`, `buffalo`, `like`}, `new` and `york` would not be included; so the signature m ith $\begin{bmatrix}1 & 0\end{bmatrix}$ and the context with $\begin{bmatrix}2 & 1\end{bmatrix}$.

(Again, use `st` e tokens for the signature.)

(g) (3 marks) Com `os_oneside` performs compared to `lesk_cos`. Why do you think this i r answer with examples.

(h) (1 mark) Suppo ing word counts as values for the vector elements, we instead used binary values, so that {`new`, `buffalo`, `york`} and {`buffalo`, `buffalo`, `like`} would be represented with $\begin{bmatrix}1 & 0 & 1 & 1\end{bmatrix}$ and $\begin{bmatrix}1 & 1 & 0 & 0\end{bmatrix}$, respectively. This is a vector representation of a set.

If we use `CosSim` for such vectors, how would this be related to the set intersection? (You do not need to implement this.)

Finally, let's try to incorporate modern word vectors in place of the bag-of-words–based method above. Relatively simple models such as the skip-gram model of word2vec can be trained on large amounts of unlabelled data; because of the large size of their training data, they are exposed to many more tokens and contexts. Once trained, word vectors can be extracted from the model and used to represent words for other tasks, usually bestowing substantial increases to performance. They also seem to exhibit some interesting semantic properties; recall the example discussed in class where if we take the vector for *king*, subtract the vector for *man*, add the vector for *woman*, and then ask which existing word vector is closest to the result, the answer will be the vector for *queen*. It stands to reason that incorporating these vectors might help improve the Lesk algorithm.

(i) (4 marks) In the `lesk_w2v` function, implement a variant of your `lex_cos` function where the vectors for the signature and context are constructed by taking the mean of the word2vec vectors for the words in the signature and sentence, respectively. Count each word once only; i.e., treat the signature and context as sets rather than multisets.

(Again, use your `stop_tokenize` to get the tokens for the signature.)

You can run your implementations on the evaluation set that we provide by running `python3 q1.py`. This will skip any unimplemented functions and display scores for the implemented ones. Report your scores in your written submission.

(j) (2 marks) Alter your code so that all tokens are lowercased before they are used for any of the comparisons, vector lookups, etc. How does this alter the different methods' performance? Why?

(Do not submit this lowercased version.)

程序代写代做CS编程辅导
WeChat: cstutorcs
Assignment Project Exam Help
Email: tutorcs@163.com
QQ: 749389476
https://tutorcs.com

程序代写代做 CS编程辅导

2. Word sense disambiguation with BERT (22 marks)

word2vec associates a[illegible]rd *type*. Newer models, such as ELMo, GPT, and BERT instead produce vect[illegible]*en*. Another way of saying this is that vector representations produced by t[illegible]*onditioned on context* and will differ depending on the context in which the [illegible]ler the two sentences *Let's play a game* and *Let's see a play*. word2vec will p[illegible]or for *play* in both cases even though they have different senses. This is why y[illegible] of the relevant vectors in the signature or sentence for `lesk_w2v` rather than[illegible]g only on the vector for *play*. Clearly, there's a parallel to how ambiguous words can have their senses resolved only in context.

(a) (4 marks) Is context really *necessary*? Assuming all that is available are wordforms and lemmata—no dependencies, semantic roles, parts of speech, etc.—can you give an example of a sentence where word order–invariant methods such as those you implemented for Q1 will never be able to completely disambiguate? If so, what is the more general pattern, and why is it impossible for the above methods to provide the correct sense for each ambiguous word? Be clear about your reasoning.[3]

WeChat: cstutorcs
Assignment Project Exam Help

But BERT (as is the case with other contextual models) doesn't produce the same vector for every instance of a particular word sense; the same vector will be produced only in the exact same context. So we cannot simply glance at the vector produced for an ambiguous word in order to determine its sense. We might consider using BERT vectors in a similar manner as we did the word2vec vectors for `lesk_w2v` above, taking the mean of the vectors to produce a vector representation for a word sequence; this turns out to perform worse than any of the methods above.[4]

Email: tutorcs@163.com
QQ: 749389476

Instead, suppose that we had a vector associated with every *sense*. If the sense vectors are related to corresponding word token vectors in some way, we could then compute similarities between possible sense vectors for a particular word token and its corresponding contextual word vector and then select the sense with the highest similarity. A simple method that works well is to run the model (BERT, in our case) over each of the sentences in a sense-annotated training set. We can gather the word tokens associated with all occurrences of a sense and take their average to represent that sense.

https://tutorcs.com

(b) (10 marks) Implement `gather_sense_vectors` in `q2.py` to assign sense vectors as described above.

(c) (2 marks) In the docstring for `gather_sense_vectors`, we point out that sorting the corpus by length before batching is much faster than leaving it as-is. Explain why this is the case.[5]

(d) (4 marks) Implement `bert_1nn` in `q2.py` to predict the sense for a word in a sentence given sense vectors produced by `gather_sense_vectors`. Keep in mind the note in the docstring about loop usage.

---

[3] Hint: re-read the preceding paragraph with this question in mind.
[4] Why?
[5] Hint: think about padding.

For this assignment, you will be using the BERT model, though the direct calls to it have been implemented for you in given helper functions. Pay attention to the notes, especially in q2.py, as they will help with and fully specify proper usage.

Finally, beware that the version of word sense disambiguation in this assignment is restricted compared to what wo[illegible]or the general case of disambiguating ambiguous words in arbitrary sentences [illegible]le you've written assumes that it is known which words need to be disambigu[illegible]ambiguous words in the sentence has already been done for you.

(e) (2 marks) Thin[illegible]r issue that would come up when attempting to use the code for this as[illegible]guate arbitrary sentences. You may consider either the Lesk variants from Q1 or the BERT-based method here (or both).

程序代写代做CS编程辅导

WeChat: cstutorcs

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

https://tutorcs.com

Notes

程序代写代做 CS编程辅导

- Running the co[illegible]t take too long, but if you would like to iterate faster, caching tokens [illegible]ntexts can be an effective speedup.
- The BERT cod[illegible]e slow and memory-intensive, so it's recommended that you run it on a [illegible]ould run on CPU in less than an hour or so, so it's definitely viable to do so if needed or preferred.) The code we've provided for this assignment will automatically use the GPU if it's available; you should never have to specify the device of a `Tensor` or call `.clone().detach()` on one. Similar to A1, you can run your Q2 code on a GPU-equipped machine via ssh to `teach.cs.toronto.edu`; use the `gpu-run-q2.sh` script to do so. As a reminder, the GPU machines are shared resources, so there are limits on how long your code is allowed to run.
- **Do not** alter the function signatures (i.e., the names of the functions or anything about the parameters that they take), and **do** make sure that your implementations of the functions return what the functions are specified to return (i.e., no extra return objects). Also, make sure your submitted files are importable in Python (i.e., make sure that `'import q1'` doesn't yield errors, and likewise for the other questions/files).
- **Do not** add any extra package or module imports. If you really would like to use a certain package or module, ask on the Piazza forum with a brief justification.
- You may add any helper functions, global variables, etc. that you wish to the files for each question (`q0.py`, `q1.py`, etc.), but do not allow your code to become overly convoluted and make sure it is still readable so that we can trace and understand your logic. You will not be submitting the other Python files, but avoid making changes to them, as your code's behaviour will be evaluated against those files as they are in the starter code.

WeChat: cstutorcs
Assignment Project Exam Help
Email: tutorcs@163.com
QQ: 749389476
https://tutorcs.com

What to submit

程序代写代做 CS编程辅导

Submission is electr[illegible] from any Teaching Labs machine using the `submit` command:

```
$ submit -c <cour[illegible]me-1> ... <filename-n>
```

where `<course>` is c[illegible] depending on which course you're registered in, and `<filename-1>` to `<fi[illegible]` files you are submitting. The files you are to submit are as follows:

- `a2written.pdf`: a PDF document containing your answers as applicable for each question. This PDF must also include a typed copy of the Student Conduct declaration as on the last page of this assignment handout. Type the declaration as it is and sign it by typing your name.
- `q0.py`: the (entire) `q0.py` file with your implementations filled in.
- `q1.py`: the (entire) `q1.py` file with your implementations filled in. Again, do not include the alterations that you implement for question 1h.
- `q2.py`: the (entire) `q2.py` file with your implementations filled in.

WeChat: cstutorcs
Assignment Project Exam Help
Email: tutorcs@163.com
QQ: 749389476
https://tutorcs.com

程序代写代做CS编程辅导

CSC 485H/2501H, Fall 2022: Assignment 2

Family name: _____

Given name: _____

Student #: _____

Date: _____

WeChat: cstutorcs

Assignment Project Exam Help

I declare that this assignment, both my paper and electronic submissions, is my own work, and is in accordance with the University of Toronto Code of Behaviour on Academic Matters and the Code of Student Conduct.

Email: tutorcs@163.com

Signature: _____

QQ: 749389476

https://tutorcs.com